

Project: Dogs, Fried Chicken or Blueberry Muffins?

Team #4

Summary:

In this project, we created a classifier for images of puppies, fried chickens and blueberry muffins.

Install Packages

```
# packages.used=c("gbm", "caret", "DMwR", "nnet", "randomForest", "EBImage", "e1071", "xgboost")
#
# # check packages that need to be installed.
# packages.needed=setdiff(packages.used,
#                          intersect(installed.packages()[,1],
#                                    packages.used))
# # install additional packages
# if(length(packages.needed)>0){
#   install.packages(packages.needed, dependencies = TRUE)
# }
```

Read in SIFT feature data

```
sift_train0 <- read.csv("./data/sift_train.csv", header=F)
label_train0 <- read.csv("./data/label_train.csv", header=F)
source("./lib/eco2121_train_gbm_baseline.r")
source("./lib/pca_features.r")
#source("./lib/new_xgboost_sift_pca100.r")

sift <- sift_train0[, -1]
```

Use PCA to reduce dimension

```
set.seed(500)
# data <- pca_features(sift, 100)
#
# # selected data with labels
# pca_train_data <- cbind(data, label_train0[,2])
# colnames(pca_train_data)[ncol(pca_train_data)] <- "label"
# pca_train_data<-as.data.frame(pca_train_data)

sift_pca<-read.csv("./data/feature_pca100.csv",header = T, as.is = T)
label<-read.csv("./data/label_train.csv",header = T,as.is = T)
dat<-cbind(label[,2],sift_pca[, -1])
colnames(dat)[1]<- "label"
```

Train and Validate set

```
set.seed(500)
# Train and test split
train_index<-sample(1:nrow(dat),0.7*nrow(dat))

xgb_variables<-as.matrix(dat[,-1]) # Full dataset
xgb_label<-dat[,1] # Full label

# Split train data
xgb_train<-xgb_variables[train_index,]
train_label<-xgb_label[train_index]
train_matrix<-xgb.DMatrix(data = xgb_train, label=train_label)

# Split test data
xgb_test<-xgb_variables[-train_index,]
test_label<-xgb_label[-train_index]
test_matrix<-xgb.DMatrix(data = xgb_test, label=test_label)
```

Baseline Model: GBM + SIFT

```
sift_train = read.csv("./data/sift_train.csv")
label = read.csv("./data/label_train.csv")
data = data.frame(label[,2], sift_train[,2:ncol(sift_train)])
colnames(data)[1] = "label"

set.seed(123)
index = sample(1:nrow(data), size=0.7*nrow(data))
train_data = data[index,]
test_data = data[-index,]

## To run the baseline model uncomment the following ##

# dat_train = training features
# label_train = labels
# K = number of folds
# d = a certain interaction depth
# system.time(result<-gbm_train(train_data[,2:ncol(train_data)],train_data$label))
# result

#
```

Our Model: XGBoost + PCA + SIFT

```
# Basic model
basic = xgboost(data = train_matrix,
                max.depth=3,eta=0.01,nthread=2,nround=50,
                objective = "multi:softprob",
                eval_metric = "mlogloss",
                num_class = 3,
                verbose = F)
```

```

# Tune the model
xgb_params_3 = list(objective="multi:softprob",
                     eta = 0.01,
                     max.depth = 3,
                     eval_metric = "mlogloss",
                     num_class = 3)

# fit the model with arbitrary parameters
xgb_3 = xgboost(data = train_matrix,
                params = xgb_params_3,
                nrounds = 100,
                verbose = F)

# cross validation
xgb_cv_3 = xgb.cv(params = xgb_params_3,
                  data = train_matrix,
                  nrounds = 100,
                  nfold = 5,
                  showsd = T,
                  stratified = T,
                  verbose = F,
                  prediction = T)

# set up the cross validated hyper-parameter search
xgb_grid_3 = expand.grid(nrounds=c(100,250,500),
                        eta = c(1,0.1,0.01),
                        max_depth = c(2,4,6,8,10),
                        gamma=1,
                        colsample_bytree=0.5,
                        min_child_weight=2,
                        subsample = 1)

# pack the training control parameters
xgb_trcontrol_3 = trainControl(method = "cv",
                              number = 5,
                              verboseIter = T,
                              returnData = F,
                              returnResamp = "all",
                              allowParallel = T)

# train the model for each parameter combination in the grid

ptm <- proc.time() ## start the time

xgb_train_3 = train(x=train_matrix, y=train_label,
                   trControl = xgb_trcontrol_3,
                   tuneGrid = xgb_grid_3,
                   method = "xgbTree")

ptm2 <- proc.time()
ptm2- ptm ## stop the clock1

##      user  system elapsed
## 545.75   83.74   397.39

```

```
# ## Time for training: 350.92s
```

```
#
```

```
head(xgb_train_3$results[with(xgb_train_3$results,order(RMSE)),],5)
```

```
##      eta max_depth gamma colsample_bytree min_child_weight subsample
## 21 0.10          4      1              0.5              2         1
## 20 0.10          4      1              0.5              2         1
## 9  0.01          6      1              0.5              2         1
## 19 0.10          4      1              0.5              2         1
## 24 0.10          6      1              0.5              2         1
##      nrounds      RMSE  Rsquared      MAE      RMSESD RsquaredSD
## 21      500 0.5110378 0.6141201 0.4175375 0.009210265 0.01961039
## 20      250 0.5116581 0.6133756 0.4187183 0.009530715 0.01967853
## 9       500 0.5134024 0.6246922 0.4314855 0.010377619 0.02250139
## 19      100 0.5137137 0.6108225 0.4225101 0.010866963 0.02121169
## 24      500 0.5138771 0.6147863 0.4187986 0.016686129 0.02848466
##      MAESD
## 21 0.007929596
## 20 0.007746908
## 9  0.004287987
## 19 0.009226362
## 24 0.010063041
```

```
# get the best model's parameters
```

```
xgb_train_3$bestTune
```

```
##      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 21      500          4 0.1      1              0.5              2         1
```

```
# # best model
```

```
bst = xgboost(data=train_matrix,max_depth=4,eta=0.1,nthread=2,nround=250,colsample_bytree=0.5,min_child_weight=1)
```

```
pred = predict(bst, test_matrix)
```

```
prediction<-matrix(pred,nrow = 3,ncol = length(pred)/3) %>%
```

```
  t() %>%
```

```
  data.frame() %>%
```

```
  mutate(label=test_label+1,max_prob=max.col(., "last"))
```

```
# ## confusion matrix of test set
```

```
confusionMatrix(factor(prediction$label),factor(prediction$max_prob),mode = "everything")
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      Reference
```

```
## Prediction  1   2   3
```

```
##      1 294   5   7
```

```
##      2  13 243  52
```

```
##      3   19  62 205
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##      Accuracy : 0.8244
```

```
##      95% CI : (0.798, 0.8488)
```

```
##      No Information Rate : 0.3622
```

```
##      P-Value [Acc > NIR] : < 2e-16
```

```

##
##           Kappa : 0.7363
## McNemar's Test P-Value : 0.01881
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3
## Sensitivity      0.9018  0.7839  0.7765
## Specificity      0.9791  0.8898  0.8726
## Pos Pred Value   0.9608  0.7890  0.7168
## Neg Pred Value   0.9461  0.8868  0.9039
## Precision        0.9608  0.7890  0.7168
## Recall           0.9018  0.7839  0.7765
## F1               0.9304  0.7864  0.7455
## Prevalence       0.3622  0.3444  0.2933
## Detection Rate   0.3267  0.2700  0.2278
## Detection Prevalence 0.3400  0.3422  0.3178
## Balanced Accuracy 0.9405  0.8369  0.8246
# ## Accuracy: 82.67%
# ## Parameters: max.depth=4, eta=0.1, nthread=2, nround=250

```